




deeplearning.ai

# Object Detection

---

## Bounding box predictions

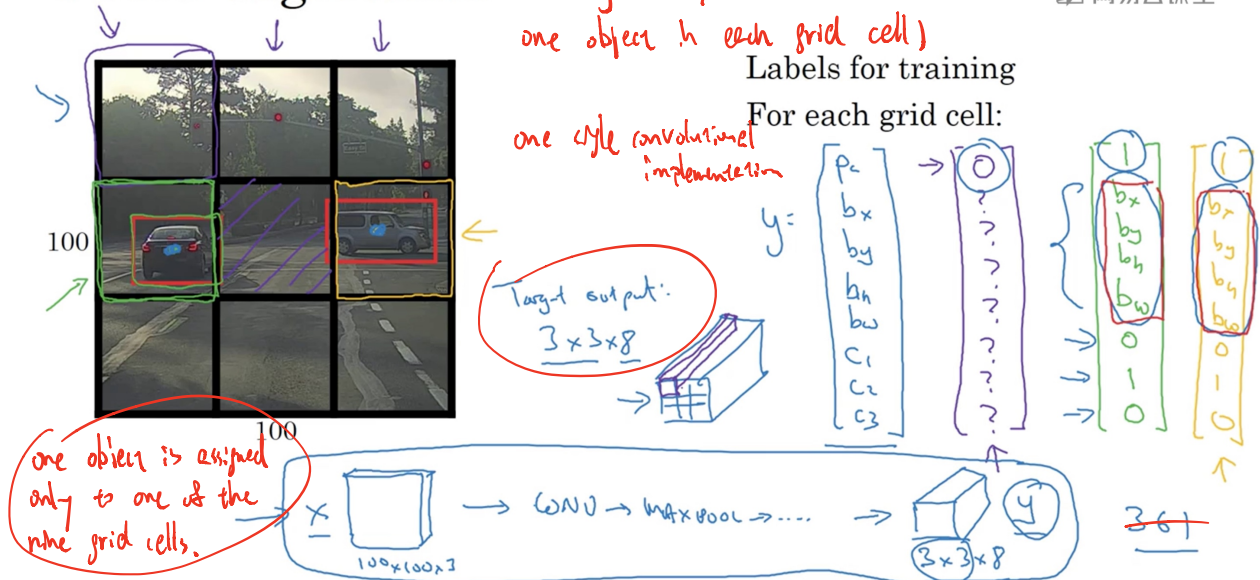
Output accurate bounding boxes



*the preferred bounding box  
isn't even quite square*

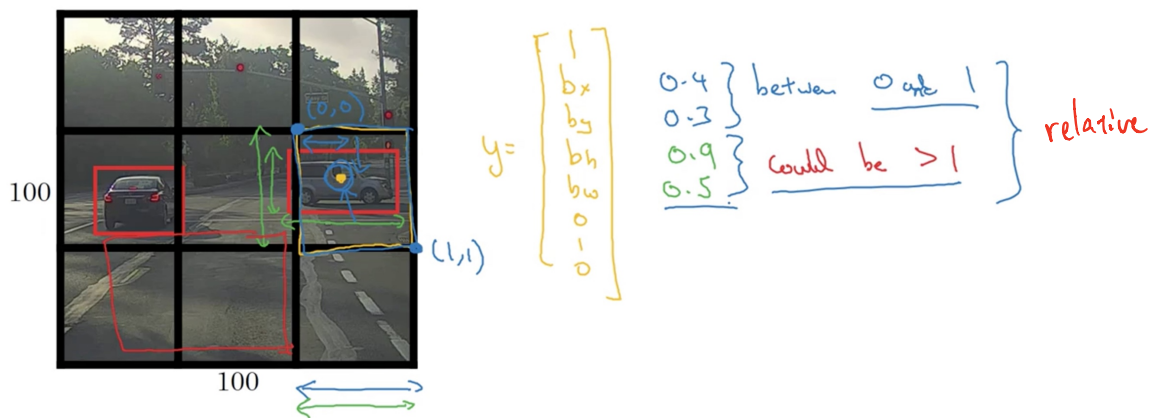


# YOLO algorithm (as long as you don't have more than one object in each grid cell)



[Redmon et al., 2015, You Only Look Once: Unified real-time object detection] → 19x19x8 (in practice) Andrew Ng

## Specify the bounding boxes



[Redmon et al., 2015, You Only Look Once: Unified real-time object detection]

Andrew Ng

