deeplearning.ai

Optimization Algorithms

Understanding mini-batch gradient descent

Training with mini batch gradient descent

Batch gradient descent

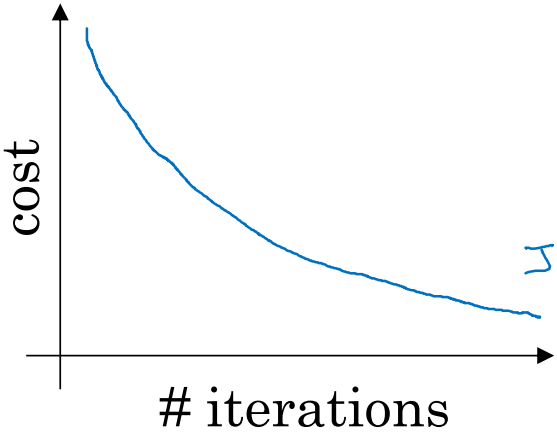

Mini-batch gradient descent

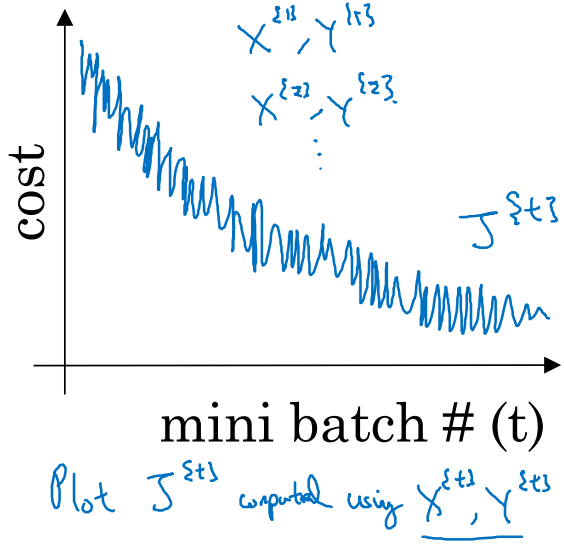

Choosing your mini-batch size

→ If mini-batch size = m : Batch gradient descent. $(X^{\{1\}}, Y^{\{1\}}) = (X, Y)$.

→ If mini-batch size = 1 : Stochastic gradient descent. Every example is it own mini-batch. $(X^{\{1\}}, Y^{\{1\}}) = (x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$

→ In practice: Somewhere in-between 1 and m

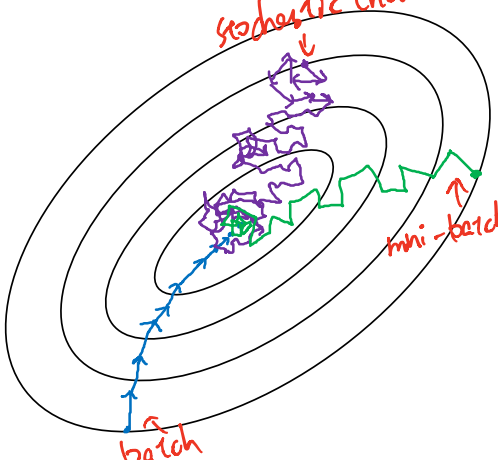

Stochastic gradient descent

Lose speedup from vectorization

In-between (mini-batch size not too big/small)

Fastest learning.
- Vectorization. (n large)
- Make progress without processing entire training set.

Batch gradient descent (mini-batch size = m)

Too long per iteration

Choosing your mini-batch size

If small train set: Use batch gradient descent.
($m \leq 2000$)

Typical mini-batch sizes:

→ 64, 128, 256, 512 (2^6, 2^7, 2^8, 2^9) — common — 2^n — 1024 (2^{10})

Make sure mini-batch fit in CPU/GPU memory.

$X^{\{t\}}, Y^{\{t\}}$